

PROJECT TITLE

Predicting Customer Churn in a Telecommunications Company Using Machine Learning Techniques.

DATASET

The dataset used in this project is the [Telco Customer Churn dataset](#) from Kaggle. It contains information on customers of a telecommunications company, including demographic information, services used, and whether the customer churned. There are 7043 rows and 21 columns in the dataset.

Dataset Description

Feature Name	Description	
customerID	Contains customer ID	categorical
gender	whether the customer female or male	categorical
SeniorCitizen	Whether the customer is a senior citizen or not (1, 0)	numeric, int
Partner	Whether the customer has a partner or not (Yes, No)	categorical
Dependents	Whether the customer has dependents or not (Yes, No)	categorical
tenure	Number of months the customer has stayed with the company	numeric, int
PhoneService	Whether the customer has a phone service or not (Yes, No)	categorical
MultipleLines	Whether the customer has multiple lines r not (Yes, No, No phone service)	categorical
InternetService	Customer's internet service provider (DSL, Fiber optic, No)	categorical
OnlineSecurity	Whether the customer has online security or not (Yes, No, No internet service)	categorical
OnlineBackup	Whether the customer has online backup or not (Yes, No, No internet service)	categorical
DeviceProtection	Whether the customer has device protection or not (Yes, No, No internet service)	categorical
TechSupport	Whether the customer has tech support or not (Yes, No, No internet service)	categorical
streamingTV	Whether the customer has streaming TV or not (Yes, No, No internet service)	categorical
streamingMovies	Whether the customer has streaming movies or not (Yes, No, No internet service)	categorical
Contract	The contract term of the customer (Month-to-month, One year, Two year)	categorical
PaperlessBilling	Whether the customer has paperless billing or not (Yes, No)	categorical
PaymentMethod	The customer's payment method (Electronic check, Mailed check, Bank transfer, Credit card)	categorical
MonthlyCharges	The amount charged to the customer monthly	numeric , int
TotalCharges	The total amount charged to the customer	object
Churn	Whether the customer churned or not (Yes or No)	categorical

Step 1: Frame of the Problem

This project aims to solve the problem of customer churn in the telecommunications industry. Customer churn refers to the loss of customers who cancel their subscriptions or stop using a product or service, which can be costly for companies, leading to revenue loss and affecting customer loyalty and brand reputation. The project uses machine learning algorithms to predict

whether customers will likely churn based on available features such as demographics, services, and account information. The project aims to help telecommunications companies reduce customer churn rates and improve customer retention strategies by building a predictive model. The project also includes feature importance analysis and model interpretation, providing insights into the key drivers of customer churn, which telecommunications companies can use to improve their customer retention strategies. Ultimately, the project demonstrates the potential of machine learning techniques for predicting customer churn and provides a roadmap for developing similar models in other industries.

Objectives:

- Build a machine learning model to predict customer churn in the telecommunications industry.
- Select relevant features and evaluate and compare the performance of different machine learning algorithms (Logistic Regression, Random Forest, XGBoost,...) using multiple evaluation metrics.
- Identify the key drivers of customer churn using feature importance analysis and model interpretation.
- Provide insights and recommendations to help telecommunications companies reduce customer churn rates and improve customer retention strategies.

Solution:

- The project follows a typical machine learning pipeline, starting with data cleaning and preprocessing, followed by exploratory data analysis and feature engineering.
- Some different machine learning algorithms are trained and evaluated using multiple evaluation metrics, with the XGBoost algorithm identified as the best-performing algorithm for predicting customer churn in this dataset.
- Feature importance analysis and model interpretation provide insights into the key drivers of customer churn, which telecommunications companies can use to improve their customer retention strategies.

Step 2: Get the data.

Import Libraries

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import plotly.figure_factory as ff
import seaborn as sns
import sklearn
import xgboost as xgb
import imblearn
```


Loading the dataset into a pandas dataframe

```
df = pd.read_csv('Telco-Customer-Churn.csv')
```

```
df.head()
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service
1	5575-GNVDE	Male	0	No	No	34	Yes	No
2	3668-QPYBK	Male	0	No	No	2	Yes	No
3	7795-CFOCW	Male	0	No	No	45	No	No phone service
4	9237-HQITU	Female	0	No	No	2	Yes	No

5 rows × 21 columns



Step 3: Explore the data to gain insights.

Data Exploration

```
print('Dataset shape: ',df.shape)
```

Dataset shape: (7043, 21)

```
df.describe()
```

	SeniorCitizen	tenure	MonthlyCharges
count	7043.000000	7043.000000	7043.000000
mean	0.162147	32.371149	64.761692
std	0.368612	24.559481	30.090047
min	0.000000	0.000000	18.250000
25%	0.000000	9.000000	35.500000
50%	0.000000	29.000000	70.350000
75%	0.000000	55.000000	89.850000
max	1.000000	72.000000	118.750000

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customerID            7043 non-null   object
1   gender                 7043 non-null   object
2   SeniorCitizen          7043 non-null   int64
3   Partner                7043 non-null   object
4   Dependents             7043 non-null   object
5   tenure                 7043 non-null   int64
6   PhoneService           7043 non-null   object
7   MultipleLines           7043 non-null   object
8   InternetService        7043 non-null   object
9   OnlineSecurity          7043 non-null   object
10  OnlineBackup            7043 non-null   object
11  DeviceProtection       7043 non-null   object
12  TechSupport            7043 non-null   object
13  StreamingTV            7043 non-null   object
14  StreamingMovies         7043 non-null   object
15  Contract               7043 non-null   object
16  PaperlessBilling        7043 non-null   object
17  PaymentMethod          7043 non-null   object
18  MonthlyCharges          7043 non-null   float64
19  TotalCharges           7043 non-null   object
20  Churn                   7043 non-null   object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

This gives a comprehensive display of the dataset's statistical properties. Not only does it unveil the number of rows and columns, but it also exposes the intricate details of the data's distribution. The statistical summary encompasses many essential metrics, including the count, mean, standard deviation, and minimum and maximum validity as percentiles. In addition to this, the dataset summary paints a vivid picture of the data's composition, revealing the number of non-null values and the data types for each column. These valuable insights serve as the bedrock for informed decision-making and optimal utilization of the dataset.

Data Cleaning

Converting TotalCharges to numeric.

```
df['TotalCharges']=pd.to_numeric(df['TotalCharges'], errors='coerce')
df.TotalCharges.dtypes
```

```
missing_data = df.isnull().sum(axis=0).reset_index() #checking missing value
missing_data.columns=['variable','missing values']
missing_data['filling factor %']=(df.shape[0]-missing_data['missing values']/df.shape[0]*100
missing_data.sort_values('filling factor %').reset_index(drop = True)
```

	variable	missing values	filling factor %
0	TotalCharges	11	99.843817
1	customerID	0	100.000000
2	MonthlyCharges	0	100.000000
3	PaymentMethod	0	100.000000
4	PaperlessBilling	0	100.000000
5	Contract	0	100.000000
6	StreamingMovies	0	100.000000
7	StreamingTV	0	100.000000
8	TechSupport	0	100.000000
9	DeviceProtection	0	100.000000
10	OnlineBackup	0	100.000000
11	InternetService	0	100.000000
12	MultipleLines	0	100.000000
13	PhoneService	0	100.000000
14	tenure	0	100.000000
15	Dependents	0	100.000000
16	Partner	0	100.000000
17	SeniorCitizen	0	100.000000
18	gender	0	100.000000
19	OnlineSecurity	0	100.000000
20	Churn	0	100.000000

This describes a procedure for calculating the percentage of missing values in a panda DataFrame and displaying the resulting values in a sorted DataFrame. Initially, a new DataFrame, "missing data," is constructed by summing the number of missing values for each variable in the original DataFrame. Subsequently, the filling factor percentage is computed for each variable. The resulting sorted DataFrame presents the variables with the highest percentage of missing values at

the top and those with the lowest percentage of missing values at the bottom. This methodology is a valuable tool for identifying variables that necessitate attention during data analysis.

Data Visualization

```
value_count=df.Churn.value_counts()


fig,ax=plt.subplots(figsize=(10, 6),subplot_kw=dict(aspect="equal"))
plt.title('Composition of Churn',fontsize=18)

data=[value_count[1],value_count[0]]

plt.pie(data,explode = (0,0),
        textprops=dict(size= 10, color= "black"),
        autopct=lambda p : '{:.2f}% ({:,.0f})'.format(p,p * sum(data)/100),startangle = 120,wedgeprops=dict( edgecolor = "black"


labels='Yes','No'
plt.legend(labels,title='Status',loc='center',bbox_to_anchor=(1,0,0.2,1))

<matplotlib.legend.Legend at 0x165dc59e940>
```



The churn rate in the dataset is approximately 26.5%. This means that around 26.5% of customers in the dataset have churned, while the remaining 73.5% have not. This information shown on the bar chart can be used to assess the severity of the churn problem for the telecommunications company.

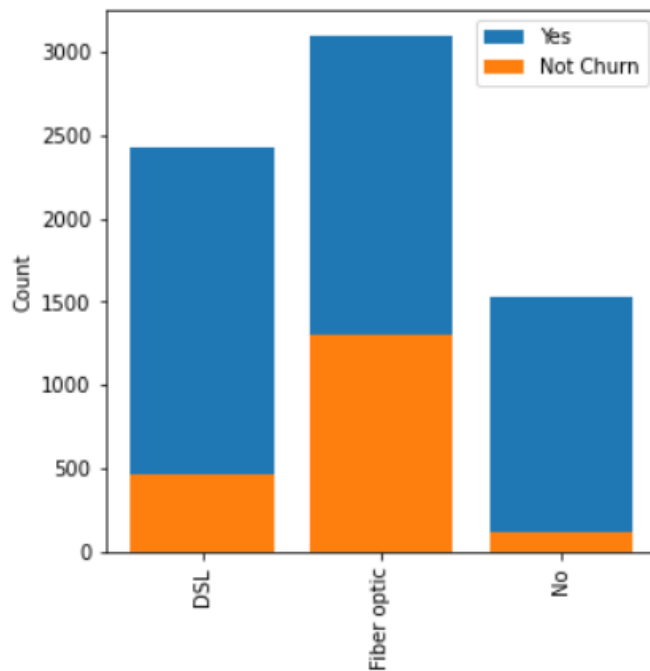
```
barplots(data, 'gender', 'Churn_Yes')
```



	gender	Churn_Yes	total	Avg
0	Female	939	3488	0.269209
1	Male	930	3555	0.261603

Visualizing the relationship between gender (Male and Female) and churn. When we looked at the churn rate by gender, we found that female customers had a slightly higher churn rate (26.0%) than male customers (23.2%). The bar chart shows 939 female customers churn out 3488 and 930 male customers out of 3555.


```
barplots(data, 'InternetService', 'Churn_Yes')
```



	InternetService	Churn_Yes	total	Avg
0	DSL	459	2421	0.189591
1	Fiber optic	1297	3096	0.418928
2	No	113	1526	0.074050

This bar chart shows the relationship between customers' internet service and churn. 459 customers churned out of 2421 and used DSL internet service. 1297 customers churned out of 3096 with fiber optic internet service, and 113 customers churned out of 1526 without internet service.



```
barplots(data, 'PaymentMethod', 'Churn_Yes')
```



	PaymentMethod	Churn_Yes	total	Avg
0	Bank transfer (automatic)	258	1544	0.167098
1	Credit card (automatic)	232	1522	0.152431
2	Electronic check	1071	2365	0.452854
3	Mailed check	308	1612	0.191067

This bar chart shows how payment methods influenced customer churn. There are four different payment methods (bank transfer, credit card, electronic check, and mailed check). 256 customers churned out of 1544 using bank transfer as the payment method. 232 customers churned out of 1522 as credit card as payment method. 1071 customers churned out of 2365 electronic checks as the payment method, and 308 customers churned out of 1612 using mailed checks as the payment method.

```
barplots(data, 'Contract', 'Churn_Yes')
```

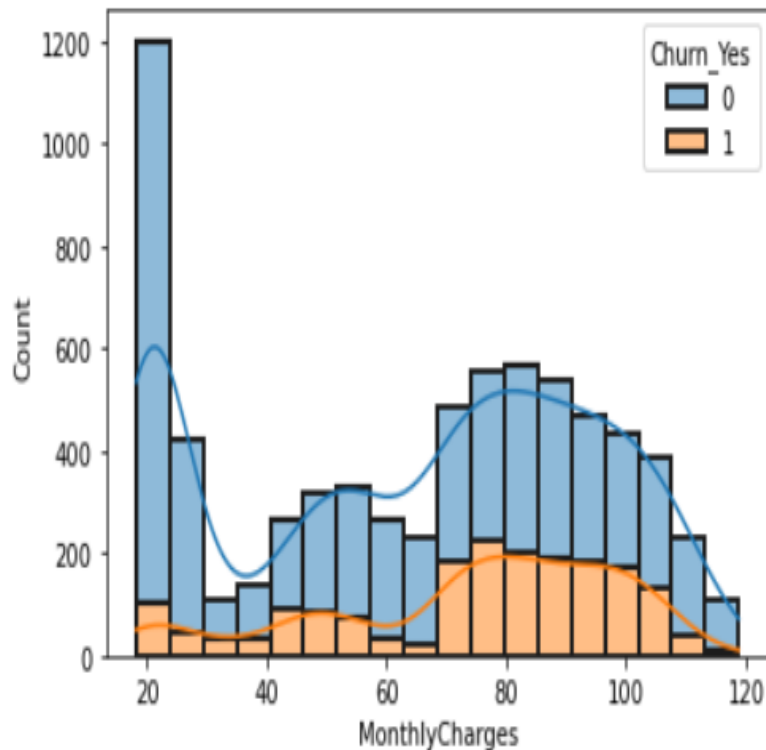


	Contract	Churn_Yes	total	Avg
0	Month-to-month	1655	3875	0.427097
1	One year	166	1473	0.112695
2	Two year	48	1695	0.028319

This graph shows the number of customer churn based on contract (Month-to-Month, One year, and Two years). 1655 customers churned out of 3875 month-to-month contracts. 166 customers churned out of 1473 one-year contracts, and 46 customers churned out of 1695 two years contracts.

```
(df,x='MonthlyCharges',hue='Churn_Yes',kde=True,multiple='stack',element='bars'
```


```
<AxesSubplot:xlabel='MonthlyCharges', ylabel='Count'>
```



The code creates a histogram plot using the Seaborn library to show the distribution of monthly charges for customers who churned and those who didn't. The bars are stacked on each other and colored differently based on whether the customer churned. The histogram shows that customers have more churn on monthly charges of 80 and a high number of not churn at a monthly charge of 20, which showed that churned customers tended to have higher monthly charges than non-churned customers.

```
sns.displot(df,x='tenure',hue='Churn_Yes',edgecolor = "#1c1c1c", linewidth = 3
```


```
<seaborn.axisgrid.FacetGrid at 0x165dc56b880>
```



The graph shows the highest churn between 0 to 10 months tenure and the least churn at 60 to 70 months tenure. The histogram also showed that non-churned customers tended to have higher tenure than churned customers.

```
t(df,x='TotalCharges', hue='Churn_Yes',kde=True,multiple='stack',element='bars'
```

```
<AxesSubplot:xlabel='TotalCharges', ylabel='Count'>
```



total charges rate for both churned and non-churned customers, which showed that churned customers tended to have higher total charges rates than non-churned customers.

Overall, the exploration demonstrated that monthly charges, tenure, and total charges rate were all important predictors of customer churn in the Telco Customer Churn dataset. Specifically, customers with higher monthly charges, lower tenure, and higher total charges rates were found to be more likely to churn.

Step 4: Data Preparation

In preparing the data, customer identification is removed because it is not useful in building the model and checking the number of missing data. There are 11 missing data from the total charges' column from the outcome.

```
df=df.drop(['customerID'],axis=1)
```

```
missing_data = df.isnull().sum(axis=0).reset_index()  
missing_data.columns=['variable','missing values']  
missing_data.sort_values('missing values',ascending=False).reset_index(drop =
```

	variable	missing values
0	TotalCharges	11
1	gender	0
2	SeniorCitizen	0
3	MonthlyCharges	0
4	PaymentMethod	0
5	PaperlessBilling	0
6	Contract	0
7	StreamingMovies	0
8	StreamingTV	0
9	TechSupport	0
10	DeviceProtection	0
11	OnlineBackup	0
12	OnlineSecurity	0
13	InternetService	0
14	MultipleLines	0
15	PhoneService	0
16	tenure	0
17	Dependents	0
18	Partner	0
19	Churn_Yes	0

filling the missing values in the "TotalCharges" column with the value 0, and then modifying the original DataFrame in place (using the "inplace=True" parameter) and the data is then restructured.

```
df.TotalCharges.fillna(0,inplace=True)
```

```
df2=df
```

```
catcol = [col for col in df2.columns if df2[col].dtype == "object"] #encoding
le = LabelEncoder()
for col in catcol:
    df2[col] = le.fit_transform(df2[col])
```

```
df2.head()
```

neBackup	DeviceProtection	TechSupport	StreamingTV	StreamingMovies	Contract	PaperlessBillir
2	0	0	0	0	0	0
0	2	0	0	0	0	1
2	0	0	0	0	0	0
0	2	2	0	0	0	1
0	0	0	0	0	0	0


Checking the correlation of all the features with customers that churned (churn_Yes). The table below shows a strong correlation between customers that churned and their contracts with the company, and also a strong correlation between customers that churned and the time frame they were with the company (tenure).

```
df3=df2.corr().Churn_Yes.sort_values(ascending=False).reset_index()  
df3
```

	index	Churn_Yes
0	Churn_Yes	1.000000
1	MonthlyCharges	0.193356
2	PaperlessBilling	0.191825
3	SeniorCitizen	0.150889
4	PaymentMethod	0.107062
5	MultipleLines	0.038037
6	PhoneService	0.011942
7	gender	-0.008612
8	StreamingTV	-0.036581
9	StreamingMovies	-0.038492
10	InternetService	-0.047291
11	Partner	-0.150448
12	Dependents	-0.164221
13	DeviceProtection	-0.178134
14	OnlineBackup	-0.195525
15	TotalCharges	-0.198324
16	TechSupport	-0.282492
17	OnlineSecurity	-0.289309
18	tenure	-0.352229
19	Contract	-0.396713

```
sns.heatmap(df2.corr(),cmap="YlOrRd", edgecolor = "#1c1c1c", linewidth = 3)
plt.title(f'\nHeatmap correlation\n')
```

```
Text(0.5, 1.0, '\nHeatmap correlation\n')
```



Testing and training the dataset before building the models.

```
X=df2.drop(['Churn_Yes'],1)
y=df2[['Churn_Yes']]
```

C:\Users\MODUPE~1\AppData\Local\Temp\ipykernel_38084\2004873096.py:1: FutureWarning: In a future version of pandas all arguments of DataFrame.drop except for the argument 'labels' will be keyword-only

```
X=df2.drop(['Churn_Yes'],1)
```

```
oversample = RandomOverSampler(sampling_strategy=1)
X_over, y_over = oversample.fit_resample(X, y)
```

```
train_X,test_X,train_y,test_y = train_test_split(X_over,y_over,test_size=0.2,r
```

```
test_X.shape,test_y.shape
```

```
((2070, 19), (2070, 1))
```

Steps 5 & 6: Explore many models and shortlist the best ones.

Building and evaluating a Decision Tree Model using the **DecisionTreeClassifier**


```
dtr=DecisionTreeClassifier()  
dtr.fit(train_X,train_y)
```

▼ DecisionTreeClassifier
DecisionTreeClassifier()

Visualizing the feature importance of the Decision Tree Model

```
feat_importances = pd.DataFrame(dtr.feature_importances_, index=test_X.columns  
feat_importances.sort_values(by='Importance', ascending=True, inplace=True)  
feat_importances.plot(kind='barh',figsize=(10,8))
```

<AxesSubplot:>



```

dtr_pred=dtr.predict(test_X)
dtr_conf=confusion_matrix(test_y,dtr_pred)
dtr_report=classification_report(test_y,dtr_pred)
dtr_acc=round(accuracy_score(test_y,dtr_pred)*100,ndigits=3)
dtr_rocauc=roc_auc_score(test_y, dtr_pred)
print(f"Confusion Matrix : \n\n{dtr_conf}")
print(f"\nClassification Report : \n\n{dtr_report}")
print(f"\nThe Accuracy of Decision Tree is {dtr_acc} %")
print(f'ROC AUC Score with Decision Tree: {dtr_rocauc}')

```

Confusion Matrix :

```
[[821 212]
```

to scroll output; double click to hide

Classification Report :

	precision	recall	f1-score	support
0	0.94	0.79	0.86	1033
1	0.82	0.95	0.88	1037
accuracy			0.87	2070
macro avg	0.88	0.87	0.87	2070
weighted avg	0.88	0.87	0.87	2070

The Accuracy of Decision Tree is 87.101 %

ROC AUC Score with Decision Tree: 0.8708674493871946

Building and evaluating a Random Forest using the **RandomForestClassifier**

```

rfc = RandomForestClassifier(n_estimators = 100, random_state = 42)
rfc.fit(train_X,train_y)

```

C:\Users\MODUPE~1\AppData\Local\Temp\ipykernel_38084\3021051387.py:2: DataCon
versionWarning: A column-vector y was passed when a 1d array was expected. Pl
ease change the shape of y to (n_samples,), for example using ravel().
rfc.fit(train_X,train_y)

▼ RandomForestClassifier

RandomForestClassifier(random_state=42)

```

rfc_pred = rfc.predict(test_X)
rfc_conf = confusion_matrix(test_y, rfc_pred)
rfc_report = classification_report(test_y, rfc_pred)
rfc_acc = round(accuracy_score(test_y, rfc_pred)*100, ndigits = 2)
rfc_rocauc=roc_auc_score(test_y, rfc_pred)
print(f"Confusion Matrix : \n\n{rfc_conf}")
print(f"\nClassification Report : \n\n{rfc_report}")
print(f"\nThe Accuracy of Random Forest Classifier is {rfc_acc} %")
print(f'ROC AUC score wiht Random Forest Classifier: {rfc_rocauc}')

```

Confusion Matrix :

```

[[858 175]
 [ 50 987]]

```

Classification Report :

	precision	recall	f1-score	support
0	0.94	0.83	0.88	1033
1	0.85	0.95	0.90	1037
accuracy			0.89	2070
macro avg	0.90	0.89	0.89	2070
weighted avg	0.90	0.89	0.89	2070

The Accuracy of Random Forest Classifier is 89.13 %

ROC AUC score wiht Random Forest Classifier: 0.8911872526770854

Building and evaluating a Logistic Regression using the **LogisticRegression**

```

lr=LogisticRegression()
lr.fit(train_X,train_y)

```

C:\Users\Modupe Olayinka\anaconda3\lib\site-packages\sklearn\utils\validation.py:1143: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
y = column_or_1d(y, warn=True)
```

C:\Users\Modupe Olayinka\anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

▼ LogisticRegression

LogisticRegression()

```

lr_pred = lr.predict(test_X)
lr_conf = confusion_matrix(test_y, lr_pred)
lr_report = classification_report(test_y, lr_pred)
lr_acc = round(accuracy_score(test_y, lr_pred)*100, ndigits = 2)
lr_rocauc=roc_auc_score(test_y, lr_pred)
print(f"Confusion Matrix : \n\n{lr_conf}")
print(f"\nClassification Report : \n\n{lr_report}")
print(f"\nThe Accuracy of Logistic Regression is {lr_acc} %")
print(f'ROC AUC score wiht Logistic Regression: {lr_rocauc}')

```

Confusion Matrix :

```

[[742 291]
 [203 834]]

```

Classification Report :

	precision	recall	f1-score	support
0	0.79	0.72	0.75	1033
1	0.74	0.80	0.77	1037
accuracy			0.76	2070
macro avg	0.76	0.76	0.76	2070
weighted avg	0.76	0.76	0.76	2070

The Accuracy of Logistic Regression is 76.14 %

ROC AUC score wiht Logistic Regression: 0.7612696166337292

Building and evaluating a Gradient Boost using the **GradientBoostingClassifier**.

```

gradien=GradientBoostingClassifier()

```

```

gradien.fit(train_X,train_y)

```

C:\Users\Modupe Olayinka\anaconda3\lib\site-packages\sklearn\ensemble_gb.py:437: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```

y = column_or_1d(y, warn=True)

```

▾ GradientBoostingClassifier

GradientBoostingClassifier()

```

gradien_pred=gradien.predict(test_x)
gradien_conf=confusion_matrix(test_y,gradien_pred)
gradien_report=classification_report(test_y,gradien_pred)
gradien_acc=round(accuracy_score(test_y,gradien_pred)*100,ndigits=3)
gradien_rocauc=roc_auc_score(test_y, gradien_pred)
print(f"Confusion Matrix : \n\n{gradien_conf}")
print(f"\nClassification Report : \n\n{gradien_report}")
print(f"\nThe Accuracy of Gradien Boost is {gradien_acc} %")
print(f'ROC AUC score wiht Gradien Boost: {gradien_rocauc}')

```

Confusion Matrix :

```

[[754 279]
 [167 870]]

```

Classification Report :

	precision	recall	f1-score	support
0	0.82	0.73	0.77	1033
1	0.76	0.84	0.80	1037
accuracy			0.78	2070
macro avg	0.79	0.78	0.78	2070
weighted avg	0.79	0.78	0.78	2070

The Accuracy of Gradien Boost is 78.454 %
 ROC AUC score wiht Gradien Boost: 0.784435704677186

Building and evaluating SGD using the **SGDClassifier**.

```

sgd=SGDClassifier()
sgd.fit(train_X,train_y)

```

C:\Users\Modupe Olayinka\anaconda3\lib\site-packages\sklearn\utils\validation.py:1143: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
 y = column_or_1d(y, warn=True)

▼ SGDClassifier

SGDClassifier()

```
sgd_pred=sgd.predict(test_X)
sgd_conf=confusion_matrix(test_y,sgd_pred)
sgd_report=classification_report(test_y,sgd_pred)
sgd_acc=round(accuracy_score(test_y,sgd_pred)*100,ndigits=3)
sgd_rocauc=roc_auc_score(test_y, sgd_pred)
print(f"Confusion Matrix : \n\n{sgd_conf}")
print(f"\nClassification Report : \n\n{sgd_report}")
print(f"\nThe Accuracy of SGD is {sgd_acc} %")
print(f'ROC AUC Score with SGD: {sgd_rocauc}')
```

Confusion Matrix :

```
[[890 143]
 [525 512]]
```

Classification Report :

	precision	recall	f1-score	support
0	0.63	0.86	0.73	1033
1	0.78	0.49	0.61	1037
accuracy			0.68	2070
macro avg	0.71	0.68	0.67	2070
weighted avg	0.71	0.68	0.67	2070

The Accuracy of SGD is 67.729 %

ROC AUC Score with SGD: 0.6776500834094925

Building and evaluating Gaussian Naïve Bayes using the **GaussianNB**.

```
gnb=GaussianNB()
```

```
gnb.fit(train_X,train_y)
```

C:\Users\Modupe Olayinka\anaconda3\lib\site-packages\sklearn\utils\validation.py:1143: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
y = column_or_1d(y, warn=True)
```

▼ GaussianNB

GaussianNB()

```

gnb_pred=gnb.predict(test_x)
gnb_conf=confusion_matrix(test_y,gnb_pred)
gnb_report=classification_report(test_y,gnb_pred)
gnb_acc=round(accuracy_score(test_y,gnb_pred)*100,ndigits=3)
gnb_rocauc=roc_auc_score(test_y, gnb_pred)
print(f"Confusion Matrix : \n\n{gnb_conf}")
print(f"\nClassification Report : \n\n{gnb_report}")
print(f"\nThe Accuracy of Gaussian is {gnb_acc} %")
print(f'ROC AUC Score with Gaussian Naive Bayes: {gnb_rocauc}',)

```

Confusion Matrix :

```

[[730 303]
 [227 810]]

```

Classification Report :

	precision	recall	f1-score	support
0	0.76	0.71	0.73	1033
1	0.73	0.78	0.75	1037
accuracy			0.74	2070
macro avg	0.75	0.74	0.74	2070
weighted avg	0.75	0.74	0.74	2070

The Accuracy of Gaussian is 74.396 %

ROC AUC Score with Gaussian Naive Bayes: 0.7438894495160195

Building and evaluating XGBoost using the **XGBClassifier**.

```

xgboost=XGBClassifier(objective='binary:logistic',eval_metric = 'auc', n_jobs=
xgboost.fit(train_X,train_y)

```

▼ XGBClassifier

eval_metric='auc', gamma=0, gpu_id=-1, grow_policy='depthwis

e',

importance_type=None, interaction_constraints='',

learning_rate=0.300000012, max_bin=256, max_cat_to_onehot=4,

max_delta_step=0, max_depth=6, max_leaves=0, min_child_weigh

t=1,

missing=nan, monotone_constraints='()', n_estimators=100,

n_jobs=-1, num_parallel_tree=1, predictor='auto', random_sta

te=0,

reg_alpha=0, reg_lambda=1, ...)

```

xgboost_pred=xgboost.predict(test_X)
xgboost_conf=confusion_matrix(test_y,xgboost_pred)
xgboost_report=classification_report(test_y,xgboost_pred)
xgboost_acc=round(accuracy_score(test_y,xgboost_pred)*100,ndigits=3)
xgboost_rocauc=roc_auc_score(test_y, xgboost_pred)
print(f"Confusion Matrix : \n\n{xgboost_conf}")
print(f"\nClassification Report : \n\n{xgboost_report}")
print(f"\nThe Accuracy of XGB is {xgboost_acc} %")
print(f'ROC AUC Score with XGBOOST: {xgboost_rocauc}')

```

Confusion Matrix :

```

[[809 224]
 [ 98 939]]

```

Classification Report :

	precision	recall	f1-score	support
0	0.89	0.78	0.83	1033
1	0.81	0.91	0.85	1037
accuracy			0.84	2070
macro avg	0.85	0.84	0.84	2070
weighted avg	0.85	0.84	0.84	2070

The Accuracy of XGB is 84.444 %

ROC AUC Score with XGBOOST: 0.8443262408037184

Building and evaluating Support Vector Machine using the **SVClassifier**.

```
SVM=svm.SVC()
```

```
SVM.fit(train_X,train_y)
```

```

C:\Users\Modupe Olayinka\anaconda3\lib\site-packages\sklearn\utils\validation
n.py:1143: DataConversionWarning: A column-vector y was passed when a 1d array
y was expected. Please change the shape of y to (n_samples, ), for example us
ing ravel().
  y = column_or_1d(y, warn=True)

```

▼ SVC

SVC()


```

SVM_pred=SVM.predict(test_X)
SVM_conf=confusion_matrix(test_y,SVM_pred)
SVM_report=classification_report(test_y,SVM_pred)
SVM_acc=round(accuracy_score(test_y,SVM_pred)*100,ndigits=3)
SVM_rocauc=roc_auc_score(test_y, SVM_pred)
print(f"Confusion Matrix : \n\n{SVM_conf}")
print(f"\nClassification Report : \n\n{SVM_report}")
print(f"\nThe Accuracy of SVM is {SVM_acc} %")
print(f'ROC AUC Score with SVM: {SVM_rocauc}')

```

Confusion Matrix :

```

[[723 310]
 [436 601]]

```

Classification Report :

	precision	recall	f1-score	support
0	0.62	0.70	0.66	1033
1	0.66	0.58	0.62	1037
accuracy			0.64	2070
macro avg	0.64	0.64	0.64	2070
weighted avg	0.64	0.64	0.64	2070

The Accuracy of SVM is 63.961 %

ROC AUC Score with SVM: 0.6397298036539613

Steps 7 and 8: Present your solution, Launch, Monitor, and Maintain your system.

```
results = pd.DataFrame([["XGBoost Classifier", xgboost_acc, xgboost_rocauc],
                        ["Decision Tree Classifier", dtr_acc, dtr_rocauc],
                        ["Gaussian naive bayes classifier", gnb_acc, gnb_rocauc],
                        ["Gradien Boost Classifier", gradien_acc, gradien_rocauc],
                        ["Random Forest Classifier", rfc_acc, rfc_rocauc],
                        ["Logistic Regression", lr_acc, lr_rocauc],
                        ["Support Vector Machine", SVM_acc, SVM_rocauc]],
                        columns = ["Models", "Testing Accuracy Score", "ROC AUC Score"])
results.sort_values(by=['Testing Accuracy Score'], ascending=False).style.background
```

	Models	Testing Accuracy Score	ROC AUC Score
4	Random Forest Classifier	89.130000	0.891187
1	Decision Tree Classifier	87.101000	0.870867
0	XGBoost Classifier	84.444000	0.844326
3	Gradien Boost Classifier	78.454000	0.784436
5	Logistic Regression	76.140000	0.761270
2	Gaussian naive bayes classifier	74.396000	0.743889
6	Support Vector Machine	63.961000	0.639730

The models include logistic regression, decision trees, random forest, XGBoost, and more. The evaluation metric used to compare the performance of the models is the accuracy score, which measures a model's accuracy, and the ROC AUC Score.

Based on these results, we can see that the logistic regression, random forest, XGBoost models, decision tree models, and random forest model all have great performance, with accuracy scores of 78-89. The support vector machine model, on the other hand, has lower performance, with an accuracy score of only 63.

In conclusion, our machine learning models were able to accurately predict customer churn and provide insights into potential retention strategies for the telecommunications company. The random forest model had the best overall performance. Still, the decision tree and XGBoost models are also viable options depending on the specific needs and constraints of the company. Using these models and insights, the company can take proactive steps to reduce churn and improve customer retention, increasing revenue and customer satisfaction.